

21CS2109AA

Operating Systems

Session 10

Memory Virtualization : Address Spaces, Memory API





Memory Virtualization

- What is **memory virtualization**?
 - OS virtualizes its physical memory.
 - OS provides an **illusion memory space** per each process.
 - It seems to be seen like **each process uses the whole memory** .



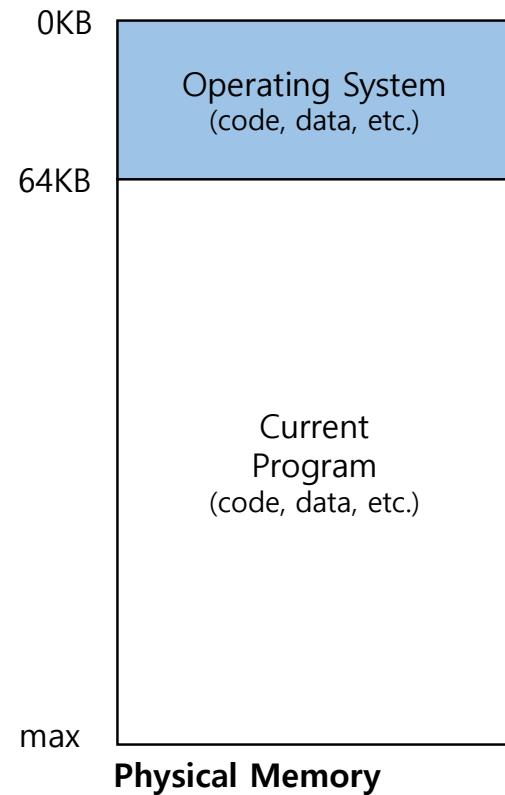
Benefit of Memory Virtualization

- Ease of use in programming
- Memory efficiency in terms of **times** and **space**
- The guarantee of isolation for processes as well as OS
 - Protection from errant accesses of other processes



OS in The Early System

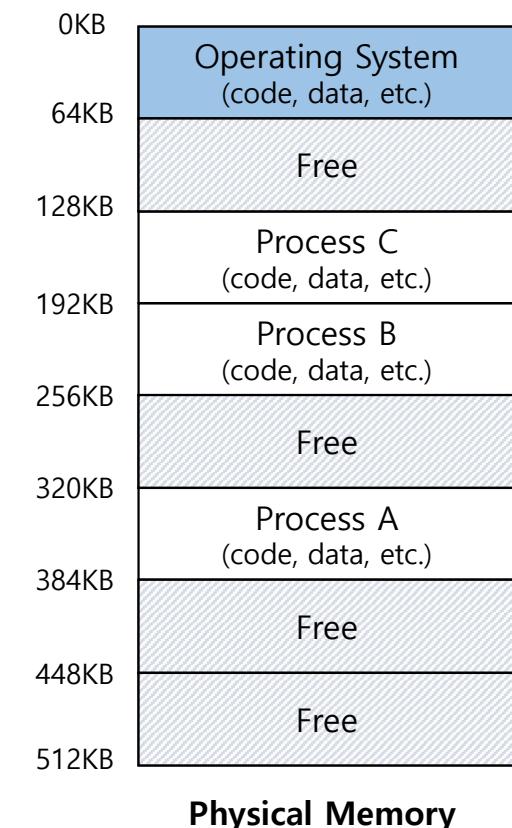
- Load only one process in memory.
 - Poor utilization and efficiency





Multiprogramming and Time Sharing

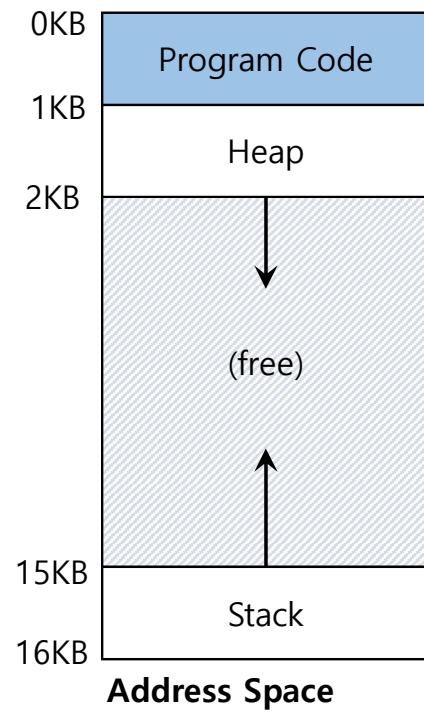
- Load multiple processes in memory.
 - Execute one for a short while.
 - Switch processes between them in memory.
 - Increase utilization and efficiency.
- Cause an important **protection issue**.
 - Errant memory accesses from other processes





Address Space

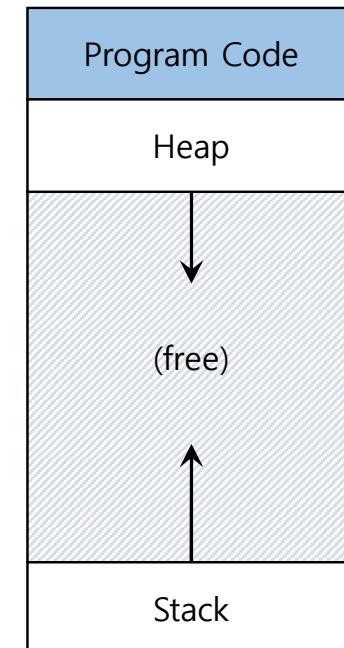
- OS creates an **abstraction** of physical memory.
 - The address space contains all about a running process.
 - That is consist of program code, heap, stack and etc.





Address Space(Cont.)

- Code
 - Where instructions live
- Heap
 - Dynamically allocate memory.
 - `malloc` in C language
 - `new` in object-oriented language
- Stack
 - Store return addresses or values.
 - Contain local variables arguments to routines.



Address Space



Virtual Address

- **Every address in a running program is virtual.**
 - OS translates the virtual address to physical address

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]){

    printf("location of code : %p\n", (void *) main);
    printf("location of heap : %p\n", (void *) malloc(1));
    int x = 3;
    printf("location of stack : %p\n", (void *) &x);

    return x;
}
```

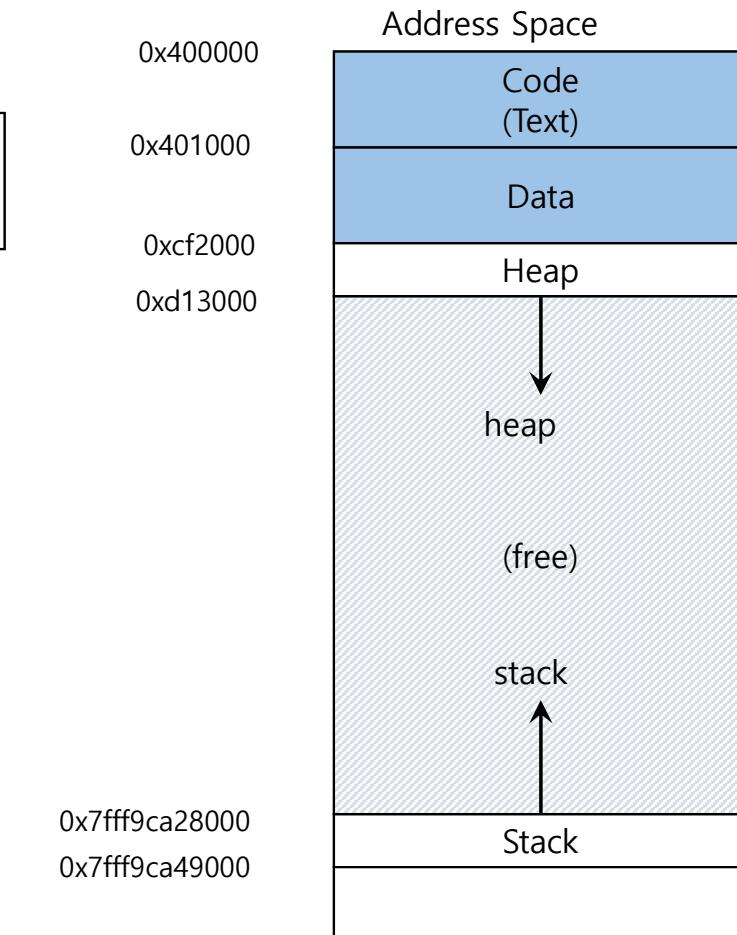
A simple program that prints out addresses



Virtual Address(Cont.)

- The output in 64-bit Linux machine

```
location of code   : 0x40057d  
location of heap   : 0xcf2010  
location of stack : 0x7fff9ca45fcc
```





Memory API: malloc()

```
#include <stdlib.h>

void* malloc(size_t size)
```

- Allocate a memory region on the heap.
 - Argument
 - `size_t size` : size of the memory block(in bytes)
 - `size_t` is an unsigned integer type.
 - Return
 - Success : a void type pointer to the memory block allocated by `malloc`
 - Fail : a null pointer



sizeof()

- Routines and macros are utilized for size in malloc instead typing in a number directly.
- Two types of results of sizeof with variables
 - The actual size of 'x' is known at run-time.

```
int *x = malloc(10 * sizeof(int));  
printf("%d\n", sizeof(x));
```

- The actual size of 'x' is known at compile-time.

```
4
```

```
int x[10];  
printf("%d\n", sizeof(x));
```

```
40
```



- When `sizeof()` is used with the data types such as `int`, `float`, `char...` etc it simply returns the amount of memory is allocated to that data types.

```
#include <stdio.h>
int main()
{
    printf("%lu\n", sizeof(char));
    printf("%lu\n", sizeof(int));
    printf("%lu\n", sizeof(float));
    printf("%lu", sizeof(double));
    return 0;
}
```

Output: 1 4 4 8



Memory API: free()

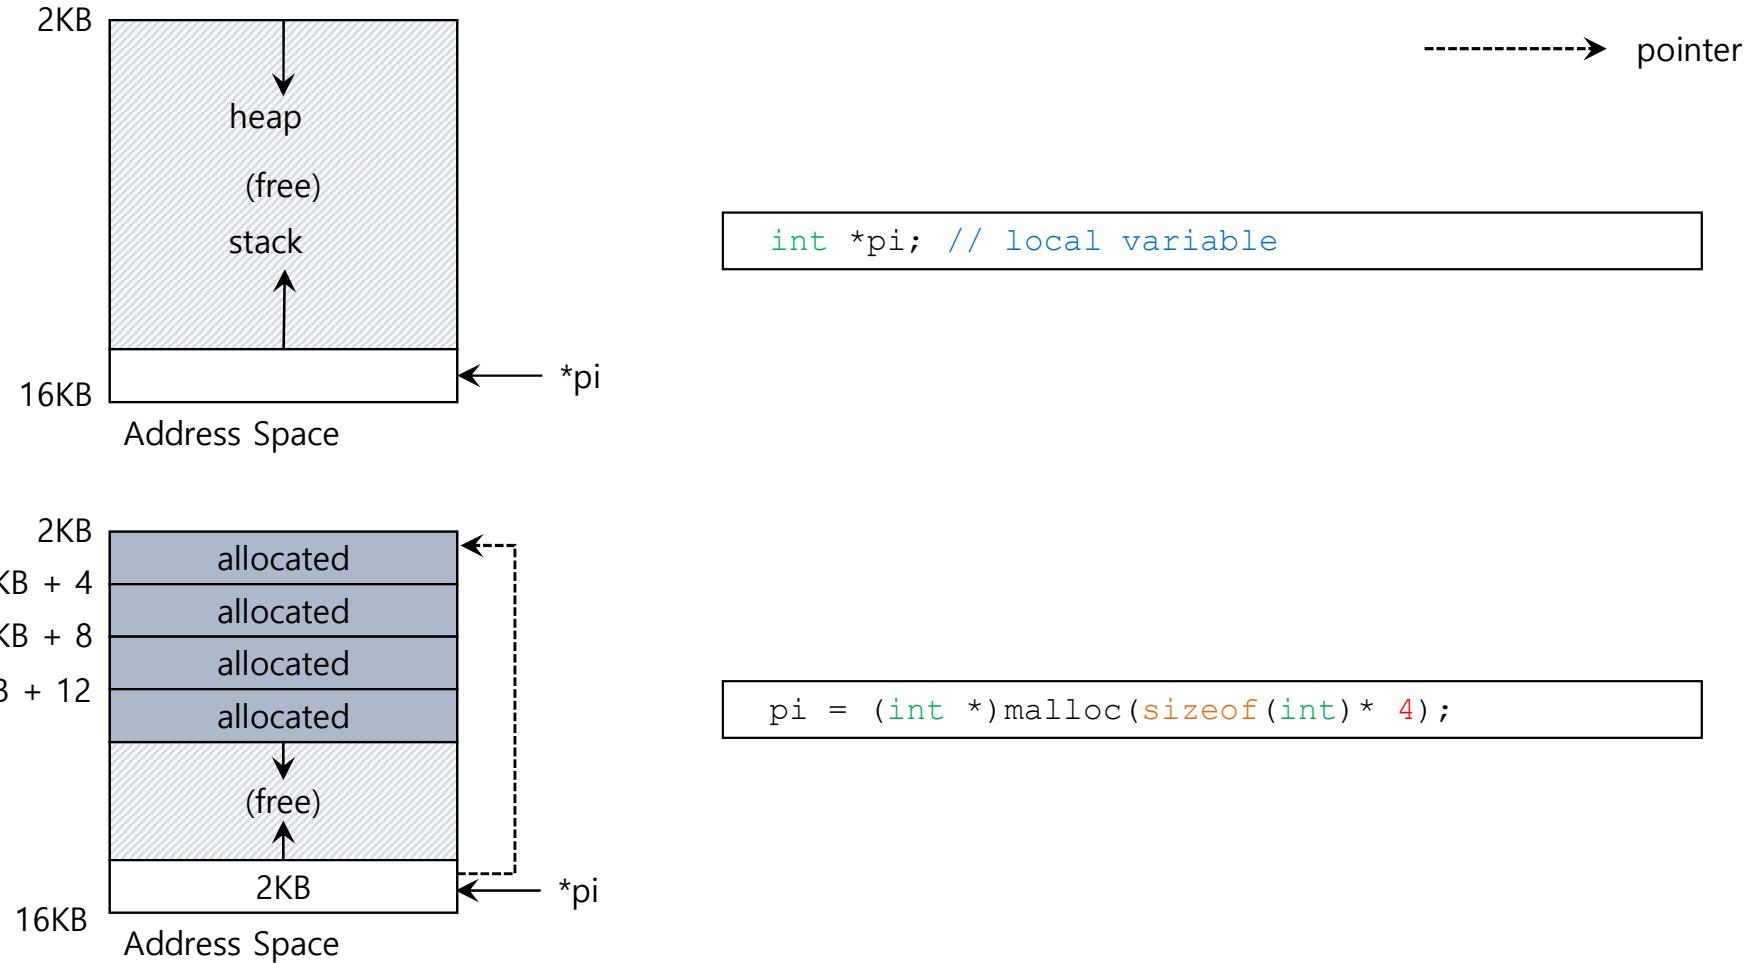
```
#include <stdlib.h>

void free(void* ptr)
```

- Free a memory region allocated by a call to malloc.
 - Argument
 - void *ptr : a pointer to a memory block allocated with malloc
 - Return
 - none

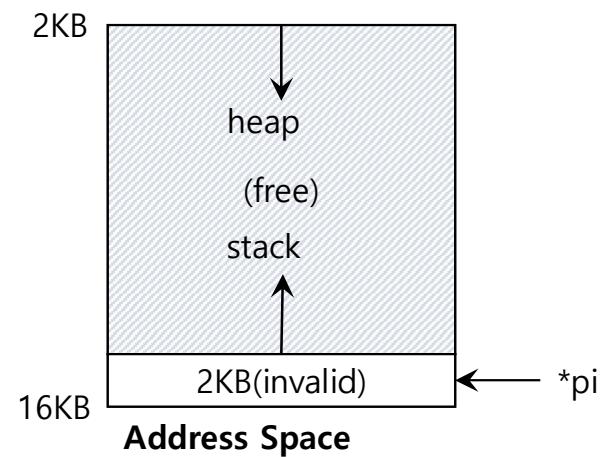
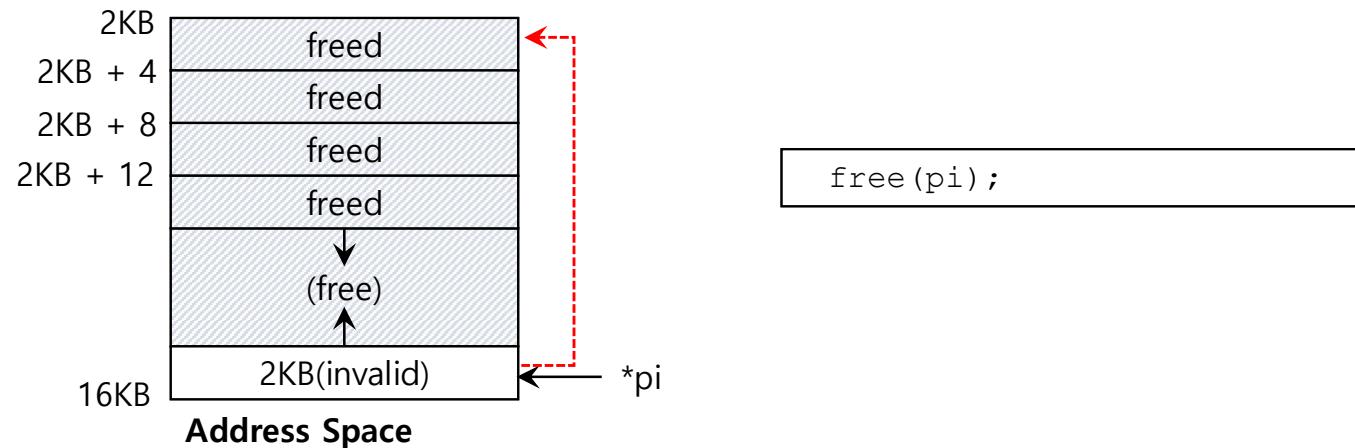


Memory Allocating





Memory Freeing

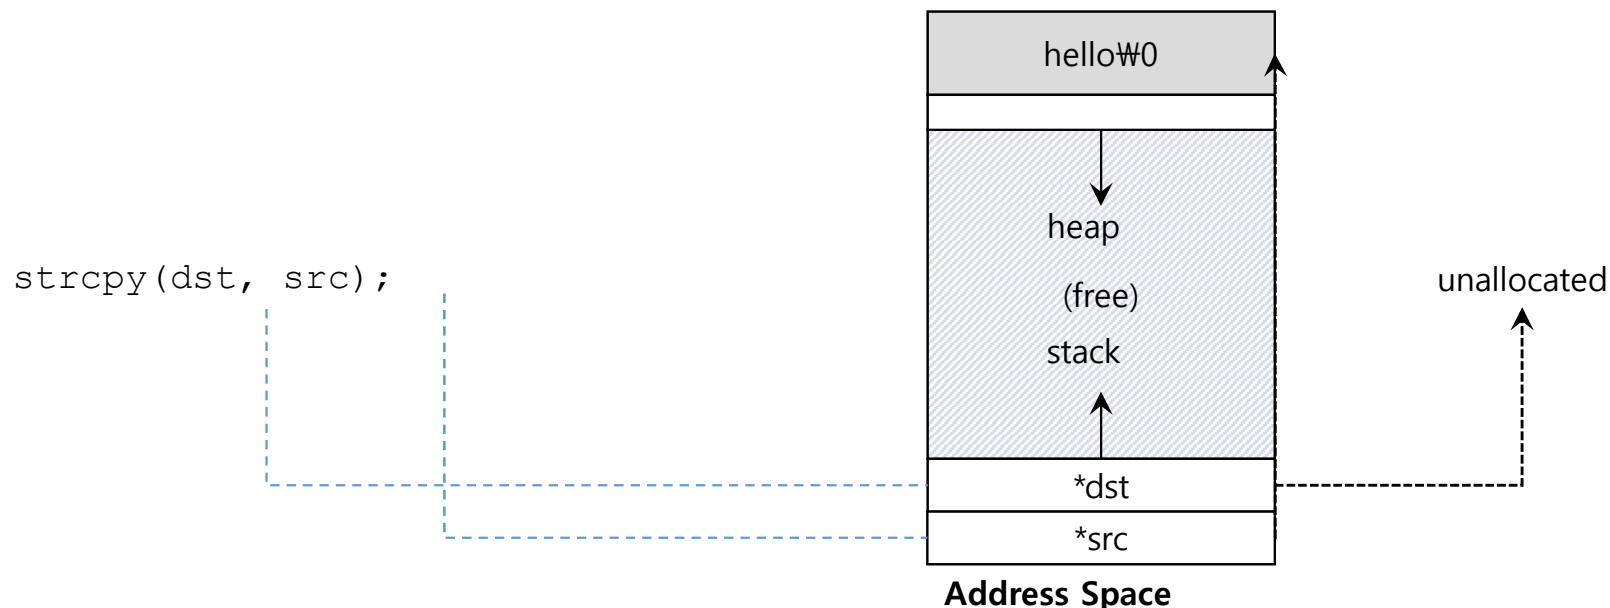




Forgetting To Allocate Memory

- Incorrect code

```
char *src = "hello"; //character string constant
char *dst;           //unallocated
strcpy(dst, src);   //segfault and die
```

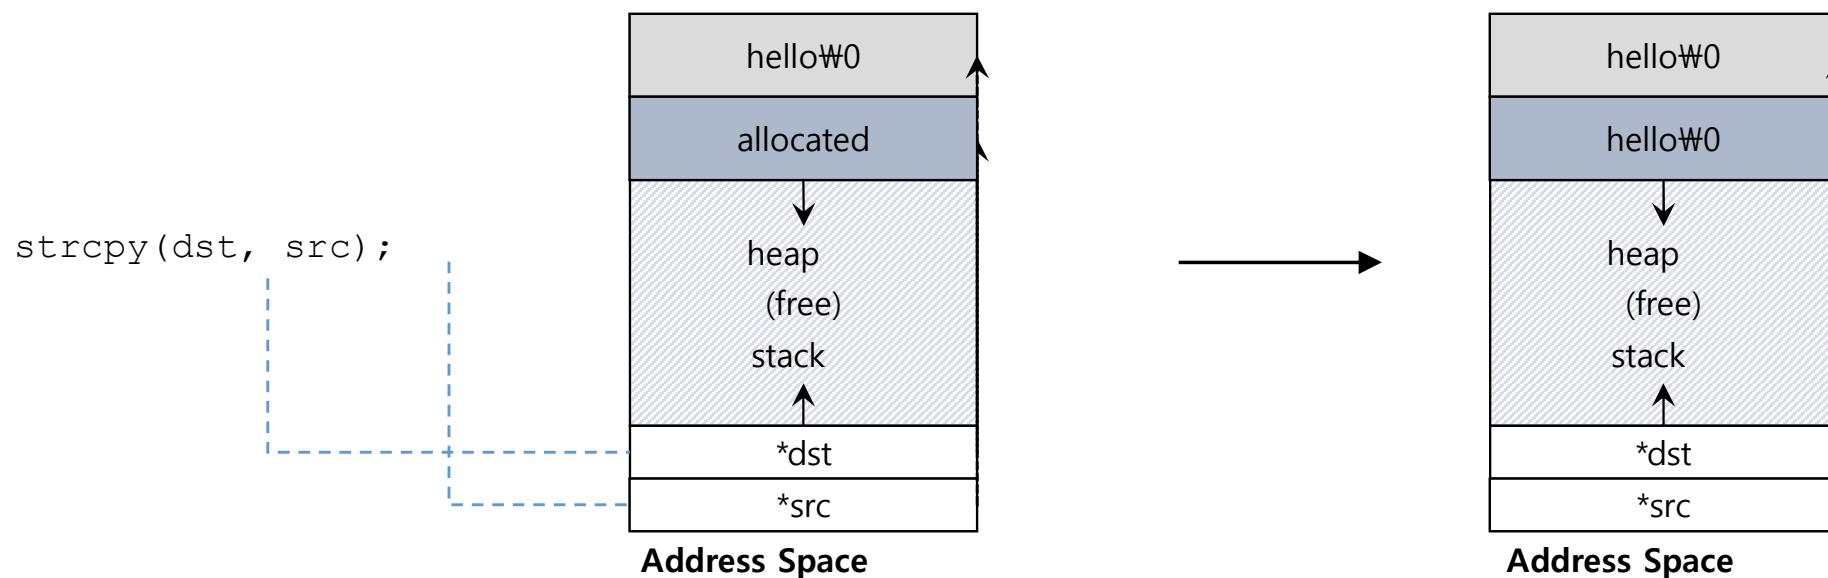




Forgetting To Allocate Memory(Cont.)

- Correct code

```
char *src = "hello"; //character string constant
char *dst (char *)malloc(strlen(src) + 1 ); // allocated
strcpy(dst, src); //work properly
```

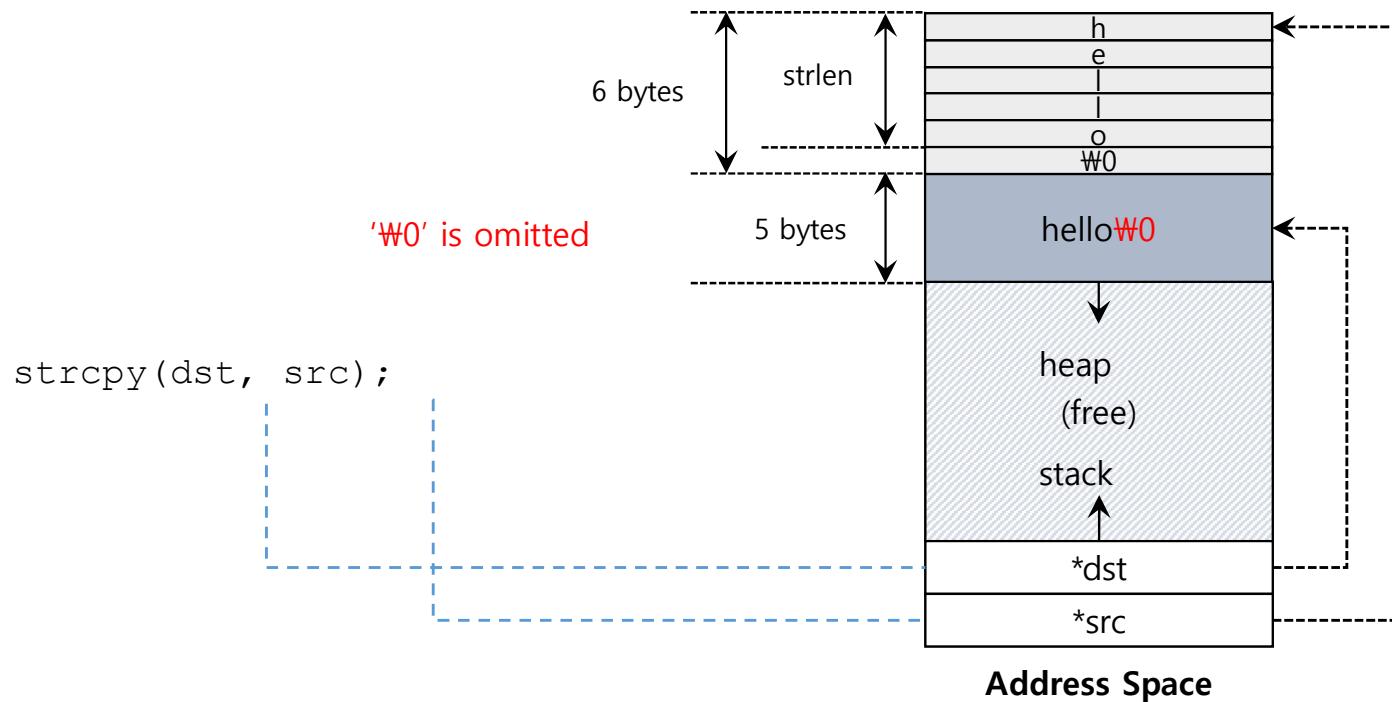




Not Allocating Enough Memory

- Incorrect code, but work properly

```
char *src = "hello"; //character string constant
char *dst (char *)malloc(strlen(src)); // too small
strcpy(dst, src); //work properly
```

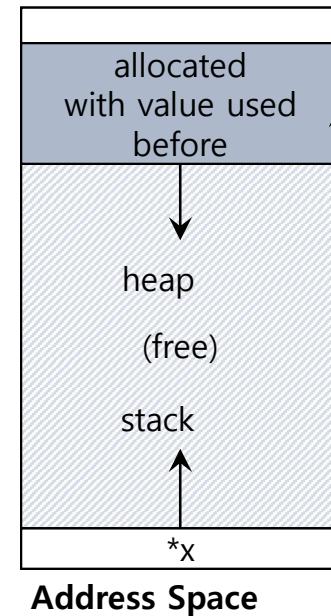
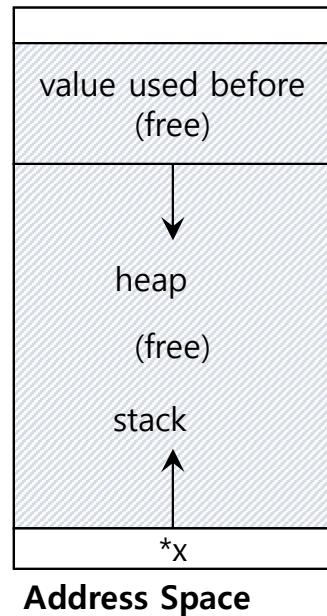




Forgetting to Initialize

- Encounter an uninitialized read

```
int *x = (int *)malloc(sizeof(int)); // allocated  
printf("*x = %d\n", *x); // uninitialized memory access
```



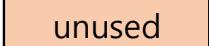


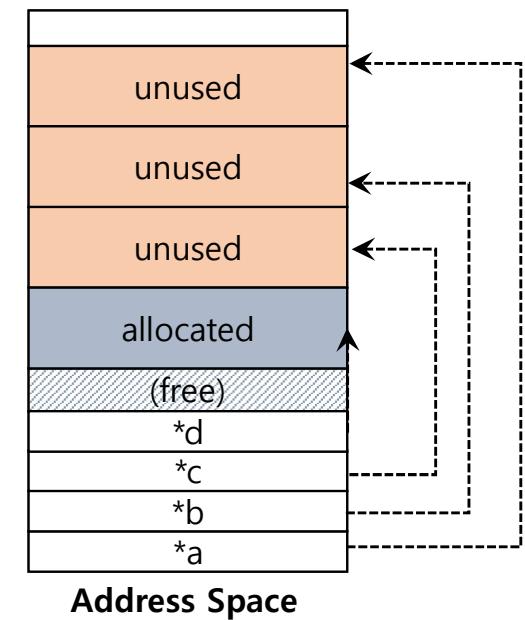
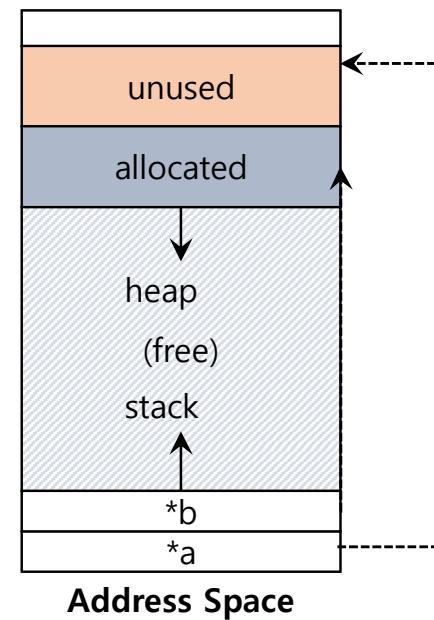
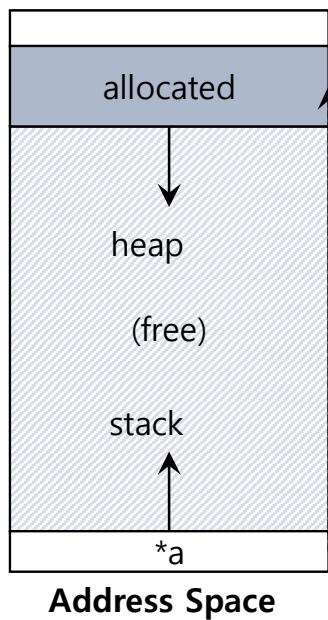
Forgetting to Initialize

- With this error, you call malloc() properly, but forget to fill in some values into your newly-allocated data type.
- Don't do this! If you do forget, your program will eventually encounter an **uninitialized read**, where it reads from the heap some data of unknown value.

Memory Leak

- A program runs out of memory and eventually dies.

 : unused, but not freed



run out of memory

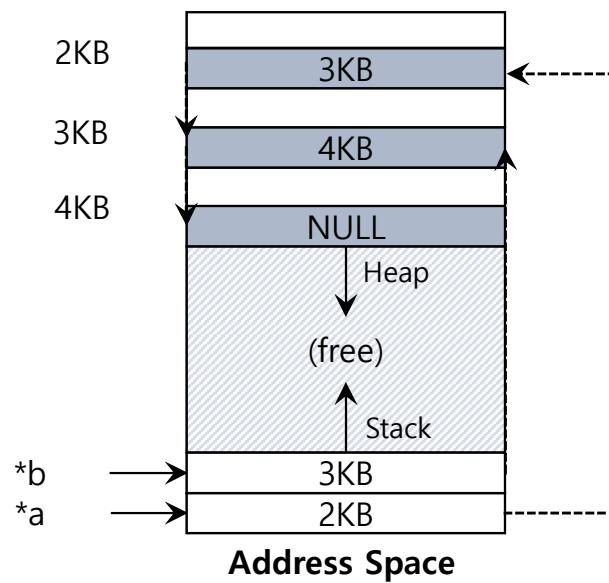
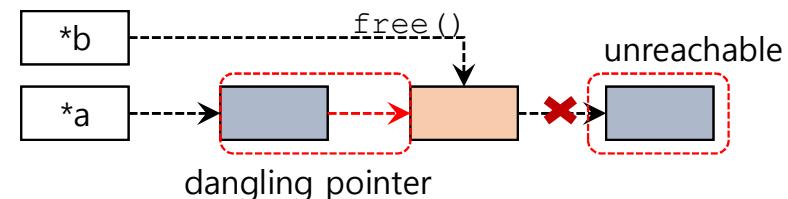
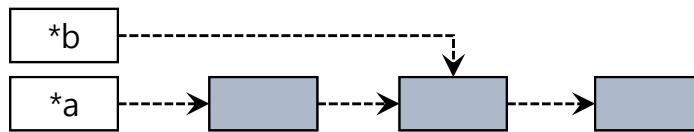


Memory Leak

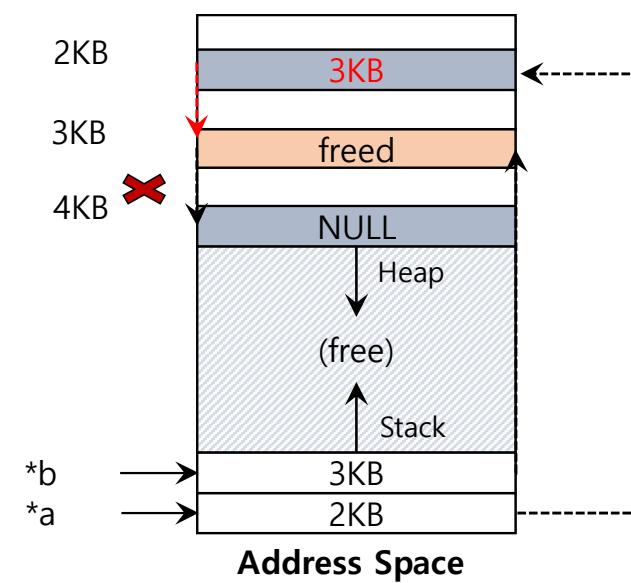
- Another common error is known as a **memory leak**, and it occurs when you forget to free memory.
- In long-running applications or systems (such as the OS itself), this is a huge problem, as slowly leaking memory eventually leads one to run out of memory, at which point a restart is required

Dangling Pointer

- Freeing memory before it is finished using
 - A program accesses to memory with an invalid pointer



`free(b)`





Dangling Pointer

- Sometimes a program will free memory before it is finished using it; such a mistake is called a **dangling pointer**, and it, as you can guess, is also a bad thing.
- The subsequent use can crash the program, or overwrite valid memory.
- e.g., you called `free()`, but then called `malloc()` again to allocate something else, which then recycles the errantly-freed memory.



Other Memory APIs: calloc()

```
#include <stdlib.h>

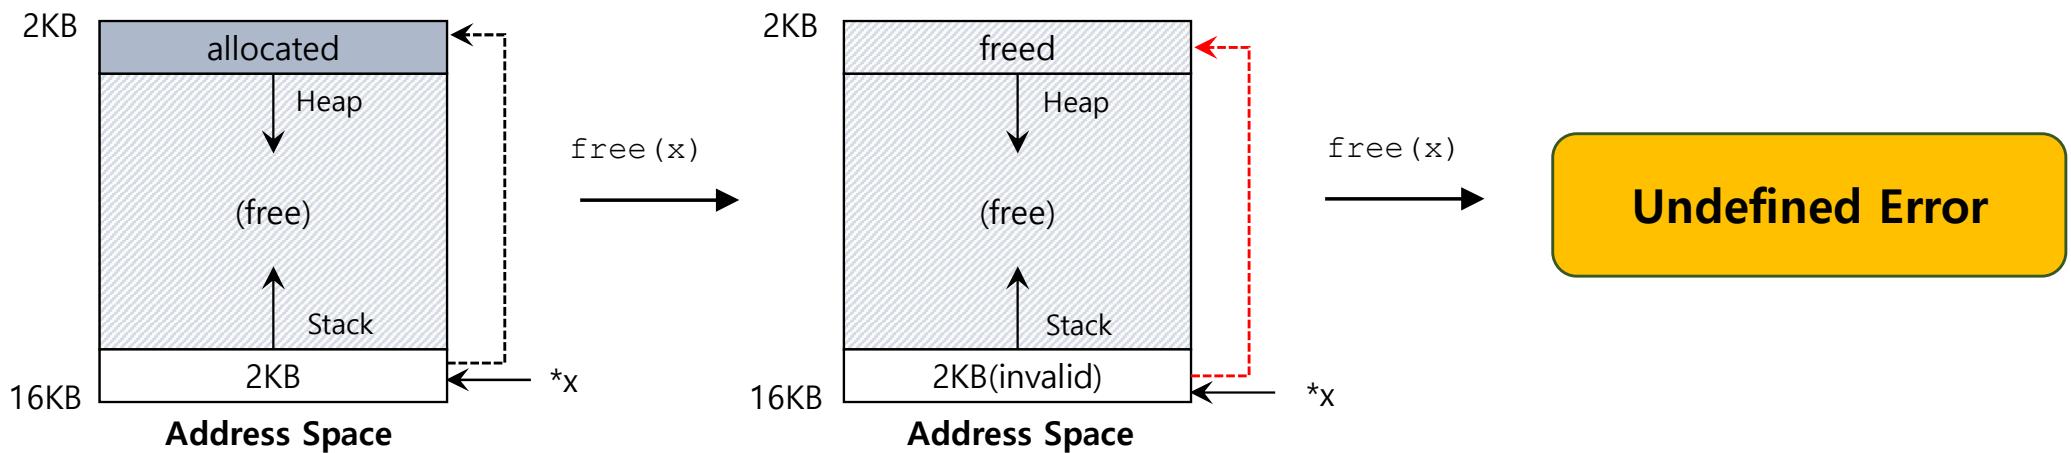
void *calloc(size_t num, size_t size)
```

- calloc() allocates the memory and also initializes the allocated memory block to zero. If we try to access the content of these blocks then we'll get 0 whereas if you want to access the content of the block which was allocated by malloc(), it will return garbage values.
- Allocate memory on the heap and zeroes it before returning.
 - Argument
 - size_t num : number of blocks to allocate
 - size_t size : size of each block(in bytes)
 - Return
 - Success : a void type pointer to the memory block allocated by calloc
 - Fail : a null pointer

Double Free

- Free memory that was freed already.

```
int *x = (int *)malloc(sizeof(int)); // allocated
free(x); // free memory
free(x); // free repeatedly
```





Double Free

- Programs also sometimes free memory more than once; this is known as the **double free**.
- **The result of doing so is undefined.**
- As you can imagine, the memory-allocation library might get confused and do all sorts of weird things; crashes are a common outcome.



Other Memory APIs: realloc()

```
#include <stdlib.h>

void *realloc(void *ptr, size_t size)
```

- The **realloc()** function is used to resize allocated memory without losing old data.
- Change the size of memory block.
 - A pointer returned by `realloc` may be either the same as `ptr` or a new.
 - Argument
 - `void *ptr`: Pointer to memory block allocated with `malloc`, `calloc` or `realloc`
 - `size_t size`: New size for the memory block(in bytes)
 - Return
 - Success: Void type pointer to the memory block
 - Fail : Null pointer



System Calls

```
#include <unistd.h>

int brk(void *addr)
void *sbrk(intptr_t increment);
```

- malloc library call use **brk** system call.
 - **brk()** and **sbrk()** change the location of the *program break*, which defines the end of the process's data segment
 - **brk()** sets the end of the data segment to the value specified by *addr*, **sbrk()** increments the program's data space by *increment* bytes.
 - Programmers **should never directly call** either **brk** or **sbrk**.



System Calls(Cont.)

```
#include <sys/mman.h>

void *mmap(void *ptr, size_t length, int port, int flags, int fd, off_t offset)
```

- mmap system call can create **an anonymous** memory region.